

Introduction to TensorFlow and Deep Learning

Lecture 1: TensorFlow basic concepts

IADS Summer School, 1st August 2022

Dr Michael Fairbank

University of Essex, UK

Email: m.fairbank@essex.ac.uk

Today's Outline

Lecture 1 (9:00am-10.30am):

- TensorFlow Basics
- Gradient Descent

Lecture 2 (11:00am-12.30pm):

- Neural Networks + Deep Learning

Lecture 3 (1:30pm-3.00pm):

- Managing Data – loading, training, regularization, saving.

Lecture 4 (3:30pm-5.00pm):

- Convolutional Neural Networks and Vision
- Introduction to Recurrent Neural Networks (* If time permits)

Course Motivations


- An introductory course
- Provide a good understanding of fundamentals of
 - TensorFlow
 - Backpropagation / AutoDiff
 - Training Neural Networks
- Not aiming to be a Keras guide to the fastest way to build neural networks
 - Hopefully give deeper understanding than that
 - However Keras will be introduced along the way

Motivations for using TensorFlow

- Ideal for neural networks and deep learning:
 - Fast tensor + matrix multiplication (using GPU if present)
 - “Automatic differentiation” is included
 - no need to write your own backprop code
 - Lots of neural network structures included
 - MLPs
 - CNNs
 - RNNs
 - LSTMs
 - Lots of optimisation algorithms included
 - Stochastic Gradient Descent
 - RMSprop
 - Adam
- It’s all supported and developed by someone else (Google)

About myself

Been developing Neural Networks since 1995



Learning algorithms for NNs and Reinforcement Learning

Recently: Applications of NNs to Electrical Power systems, Motors, Question and Answering, [Target Space](#)

Acknowledgements:

- Many pictures were created by others
- Some python code scripts were based on code from e.g. tensorflow.org

TensorFlow – Installation

- Requires Python 3.
 - Java/C/Go flavours also available, but not as stable
- Usually it's simple “pip install tensorflow” for Python
 - See <https://www.tensorflow.org/install/>
- Should have Python3 + TensorFlow v 2.9.0 installed for this course.
 - CPU version
 - Also check you have matplotlib, pandas, numpy, jupyter-notebooks:
 - pip install matplotlib pandas numpy notebook tensorflow
 - Can use google colab if you prefer

TensorFlow – Check installation

```
import tensorflow as tf
print(tf.__version__)
2.9.0
```

Try this.

(Use the accompanying jupyter notebook:

- lecture1-notebook.ipynb
- Run the appropriate notebook cell with “ctrl+enter”)

What's a tensor?

- In TensorFlow, a “rank n” tensor is equivalent to an “n dimensional array”.
- [1 2 3 4] A rank-1 tensor (a vector)
- [[1 2]
[3 4]] A rank-2 tensor (a matrix)
- 2 A rank-0 tensor (a scalar)
- Aim to build our algorithms out of tensors and matrix multiplications
 - so we can pass their manipulation off to the graphics card / low-level compiled code.

Basic concepts- Tensor scalars, and numpy

```
a=tf.constant(2,tf.float32)
b=tf.constant(3,tf.float32)
c=tf.add(a,b)
print(c)
```

A rank-0 tensor (i.e. a scalar) of type float32

```
tf.Tensor(5.0, shape=(), dtype=float32)
```

This just says c is a Tensor of shape() (rank zero), type float32, with value 5.0

```
print(c.numpy())
5.0
```

Converts from tensorflow datatype to a numpy datatype

You can try any of these in the lecture1 notebook.

Don't rush to keep up though: you can go back after the lecture and fill in any gaps

Basic concepts – tensor addition

```
a=tf.constant([[1,2],[3,4]])  
b=tf.constant([[5,6],[8,9]])  
c=tf.add(a,b)  
print(c.numpy())
```

$$\begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} + \begin{pmatrix} 5 & 6 \\ 8 & 9 \end{pmatrix} = \begin{pmatrix} ? & ? \\ ? & ? \end{pmatrix}$$

```
[[ 6  8]  
 [11 13]]
```



Numpy Array

Basic concepts – tensor multiplication

```
a=tf.constant([[1,2],[3,4]])  
b=tf.constant([[5,6],[8,9]])  
c=tf.multiply(a,b)  
print(c.numpy())
```

Elementwise multiplication
("Hadamard product")

$$c_{ij} = b_{ij}b_{ij}$$

$$\begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} * \begin{pmatrix} 5 & 6 \\ 8 & 9 \end{pmatrix} = \begin{pmatrix} ? & ? \\ ? & ? \end{pmatrix}$$

```
[[ 5  12]  
 [24  36]]
```

[https://en.wikipedia.org/wiki/Hadamard_product_\(matrices\)](https://en.wikipedia.org/wiki/Hadamard_product_(matrices))

Basic concepts – matrix multiplication

```
a=tf.constant([[1,2],[3,4]],tf.float32)
b=tf.constant([[1],[1]],tf.float32)
c=tf.matmul(a,b)
print(c.numpy())
[[3]
 [7]]
>>>
```

← A rank-2 tensor (i.e. a 2*2 matrix)

← A rank-2 tensor (a 2*1 matrix)

← Matrix multiplication
(c.f. previous slide which was
tf.multiply = elementwise
multiplication)

$$\begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} \begin{pmatrix} 1 \\ 1 \end{pmatrix} = \begin{pmatrix} ? \\ ? \end{pmatrix}$$

Basic concepts – datatypes

```
>>>a=tf.constant(3.2, tf.float32)
>>>b=tf.constant(3, tf.int32)
>>>c=tf.constant([1,2,3], tf.float32)
>>>d=tf.constant(5)
>>>e=tf.constant(5.0)
```

When you create a tensor, you should specify its datatype

← This defaults to int32

← This defaults to float32

- Also have tf.float64, tf.int64, tf.bool
 - See https://www.tensorflow.org/api_docs/python/tf/dtypes/DType

Basic concepts – casting datatypes (1)

- You can convert one datatype to another as follows:

```
a=tf.constant([[1,2],[3,-4]],tf.float32)
print(tf.cast(a,tf.int32).numpy())
[[ 1  2]
 [ 3 -4]]
```

This is now an integer tensor

```
b=tf.constant([True, False, True], tf.bool)
print(tf.cast(b,tf.int32).numpy())
[1 0 1]
```

Bools cast using True=1,
False=0

Basic concepts – casting datatypes (2)

- You can't add datatypes that don't match

```
>>>a=tf.constant(3.0, tf.float32)
>>>b=tf.constant(3, tf.int32)
>>>c=tf.add(a,b)
...
tensorflow.python.framework.errors_impl.InvalidArgumentError:
cannot compute Add as input #1(zero-based) was expected to be a
float tensor but is a int32 tensor [Op:Add]
```

- We have to cast like this:

```
c=tf.add(a,tf.cast(b,tf.float32))
print(c.numpy())
6.0
```


Basic concepts – tensor shape (1)

- Tensor shapes must match for most operations

- E.g. $\begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} + (1 \ 2 \ 3) = ???$

```
a=tf.constant([1,2])  
b=tf.constant([2,3,1])  
f=tf.add(a,b)
```

```
...
```

```
tensorflow.python.framework.errors_impl.InvalidArgumentError: Incompatible  
shapes: [2] vs. [3] [Op:Add]
```

```
...
```

Basic concepts – tensor shape (2)

- However there is a shorthand that violates these size-matching rules
- When the rank of one matrix is less than the other it tries to add them in the most sensible way (if possible)

$$\bullet \begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} + 1 = \begin{pmatrix} 2 & 3 \\ 4 & 5 \end{pmatrix}$$

```
a=tf.constant([1,2])  
b=tf.constant(1)  
print(tf.add(a,b).numpy())  
[2 3]
```

Basic concepts – tensor shape (3)

```
a=tf.constant([[1,2],[3,4]])  
b=tf.constant([10,20])  
print(tf.add(a,b).numpy())  
...
```

$$\bullet \begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} + (10 \quad 20) = \begin{pmatrix} ? & ? \\ ? & ? \end{pmatrix}$$

- This behaviour is called “broadcasting”
 - For the precise rules, see <https://docs.scipy.org/doc/numpy/user/basics.broadcasting.html>

Elementwise Tensor operations

- These elementwise operations produce a tensor of equal size to the input

```
a=tf.constant([[1,2],[3,-4]],tf.float32)
print(tf.square(a).numpy())
```

```
[[ 1.  4.]
 [ 9. 16.]
```

```
print(tf.abs(a).numpy())
```

```
[[ 1.  2.]
 [ 3.  4.]
```

```
print(tf.tanh(a).numpy())
```

```
[[ 0.76159418  0.96402758]
 [ 0.99505472 -0.99932921]]
```

Comparison Tensor operations

- See https://www.tensorflow.org/api_docs/python/tf/math/greater

```
a=tf.constant([1,2,3])  
b=tf.constant([5,1,7])  
print(tf.greater(a,b).numpy())
```

- Which elements of matrix a are bigger than those corresponding elements of b ?

```
[False  True False]
```

```
print(tf.greater(a,1).numpy())
```

```
[False  True  True]
```

This is “broadcasting” the mismatching tensor sizes, i.e. changing 1 to [1,1,1] before comparison

Basic concepts – operator shorthand

- Some math functions have shorthand operators.
- The most common ones are:

a+b tf.add(a,b)

a-b tf.subtract(a,b)

a*b tf.multiply(a,b)

a>b tf.greater(a,b)

Q: Elementwise multiplication, or matrix product?

```
a=tf.constant([[1,2],[3,4]])  
b=tf.constant([[5,6],[8,9]])  
print((a+b).numpy())
```

```
[[ 6  8]  
 [11 13]]
```

Basic concepts – variables vs. constants (1)

Unlike “constants”, all “Variables” can be updated.

```
W = tf.Variable([0.3], tf.float32)
W.assign([-1.0])
print(W.numpy())
[-1.]
```

Constants however, once created, cannot be “reassigned”

But (confusingly) you can do:

```
x = tf.constant([0.3], tf.float32)
x = tf.constant([-1], tf.float32)
print(x.numpy())
[-1.]
```

Q: What’s the difference between this reassignment and `W.assign(...)`?

Basic concepts – variables vs. constants (2)

Variables also have `assign_add` and `assign_sub`

```
W = tf.Variable([0.3], tf.float32)
W.assign_add([1.0])
print(W.numpy())
[1.3]
W.assign_sub([1.0])
print(W.numpy())
???
```

← Analogous to $W+=1$

Aggregation functions (1)

```
a=tf.constant([[1,2],[3,4]],tf.float32)
print(tf.reduce_sum(a).numpy())
10.0
print(tf.reduce_mean(a).numpy())
2.5
print(tf.reduce_max(a).numpy())
4.0
```

Aggregation functions (2)

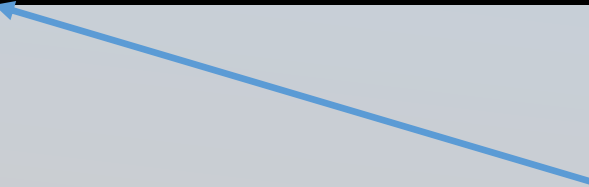
Challenge question:

```
a=tf.constant([[1,2],[3,4]],tf.float32)
print(tf.reduce_sum(tf.cast(a>1,tf.float32)).numpy())
```

Aggregation functions (3)

- Argmax counts the index at which the max element appears
- Reduce_max vs argmax:


```
a=tf.constant([4,0,5,-4],tf.float32)
print(tf.reduce_max(a).numpy())
5.0
print(tf.argmax(a).numpy())
2
```




This indicates the 3rd element
(since indices start counting at zero)

Aggregation functions across an axis (1)

- Sum the elements of a matrix across one given axis:



Axis 0




Axis 1

```
a=tf.constant([[1,2],[3,4]])
print(tf.reduce_sum(a, axis=0).numpy())
[4 6]
print(tf.reduce_sum(a, axis=1).numpy())
[3 7]
```

Aggregation functions across an axis (2)

- Aggregate elements of a matrix across one given axis:



```
a=tf.constant([[5,10,0],[3,4,12]])  
print(tf.reduce_max(a, axis=0).numpy())  
[ 5 10 12]  
print(tf.argmax(a, axis=1).numpy())  
[1 2]
```

Summary

- Tensors
 - Rank, shape, datatypes
- .numpy()
- Constants vs Variables
- Operators on tensors
 - Math functions
 - Add, multiply, matmult, tanh, greater
 - Aggregate functions
 - Max, argmax, reduce_sum

Automatic differentiation

Automatic differentiation (Autodiff) (1)

TensorFlow knows that:

$$y = x^2$$

$$\Rightarrow \frac{dy}{dx} = 2x$$

```
x = tf.Variable(3.0, tf.float32)
with tf.GradientTape() as tape:
    y=tf.pow(x,2.0)
dydx=tape.gradient(y, x)
print(dydx.numpy())
6.0
```


Automatic differentiation (Autodiff) (2)

- **Autodiff** is fast and exact differentiation
 - Not numerical differentiation (which is neither exact nor fast)
 - Not symbolic differentiation either
 - it's something in between
 - If you give it code to compute a function $f(x)$, it will write corresponding program code for you that calculated df/dx
- Further reading: <https://justindomke.wordpress.com/2009/02/17/automatic-differentiation-the-most-criminally-underused-tool-in-the-potential-machine-learning-toolbox/>

Automatic differentiation (Autodiff) (3)

Autodiff also works when there is more than one input variable:

$$f(x, y) = 3x^2 + y$$
$$\Rightarrow \frac{\partial f}{\partial x} = 6x, \frac{\partial f}{\partial y} = 1$$

“Partial derivatives”

```
x=tf.Variable(4.0,tf.float32)
y=tf.Variable(2.0,tf.float32)
with tf.GradientTape() as tape:
    f=tf.pow(x,2.0)*3.0+y
[dfdx, dfdy]=tape.gradient(f, [x,y])
print(dfdx.numpy(), dfdy.numpy())
```

← Fetching two derivatives at once

```
24.0 1.0
```

Automatic differentiation (Autodiff) (4)

If you want a derivative w.r.t. a constant, then you need `tape.watch(...)`

```
x=tf.constant(4.0,tf.float32)
with tf.GradientTape() as tape:
    tape.watch(x)
    f=tf.pow(x,3.0)
dfdx=tape.gradient(f, x)
print(dfdx.numpy(), dfdy.numpy())
```

A “constant” requires watching

```
24.0 1.0
```

Automatic differentiation (Autodiff) (5)

Autodiff also works when the input variables are higher rank tensors

$$f(W, x, y) = \|\tanh(Wx) - y\|^2$$



E.g. a 2*2 matrix

Automatic differentiation (Autodiff) (6)

- AUTODIFF makes neural-network training much easier to program
 - Autodiff replaces “backpropagation” programming
 - Backpropagation is replaced by autodiff

Introduction to Gradient Descent

Introduction to Gradient Descent





Image credit: https://sebastianraschka.com/Articles/2015_singlelayer_neurons.html

$$w_{t+1} = w_t - \eta \frac{\partial J}{\partial w}$$

Example 1D Gradient Descent problem

- Use gradient descent to find the minimum of $y = x^2 - 4x + 4$



- Start with any x_0
- Iterate with $x_{t+1} = x_t - \eta \frac{dy}{dx}$
- η is “learning rate”

Q: What will happen if η too large, too small or just right??

Example 1D Gradient Descent problem

Exercise 1: Use gradient descent to find the minimum of $y = x^2 - 4x + 4$. Complete this code in the accompanying jupyter notebook under “Exercise 1”

```
import tensorflow as tf

eta = 0.1 # learning rate
x = tf.Variable(10.0, tf.float32) # arbitrary initial value

for i in range(50):
    with tf.GradientTape() as tape:
        y=#TODO put in formula for y in terms of x here
        dydx=tape.gradient(# TODO finish this line
        x.assign(#TODO finish x_(t+1)=x_t-eta*dydx
    print("iteration:",i, "x:", x.numpy(), "y:", y.numpy())
```

Example 1D Gradient Descent problem

Observations:

- We didn't need to give the iterative variable's steps different variable names x_1, x_2, \dots . We just called them all "x"
 - So $x_{t+1} = x_t - \eta \frac{dy}{dx}$ became $x = x - \eta \frac{dy}{dx}$
- It looks pretty inefficient:
 - Recalculates out the automatic differentiation formula every step of loop!

Example 1D Gradient Descent problem

Observations:

We should have correctly found the minimum
at $x=2$, $y=0$:


iteration: 41 x : 2.00068 y : 4.76837e-07

iteration: 42 x : 2.00054 y : 4.76837e-07

iteration: 43 x : 2.00044 y : 0.0

iteration: 44 x : 2.00035 y : 0.0

...



Example 1D Gradient Descent problem

Observations:

- If eta too high, it fails
- If eta too low, solution is very slow.

Recommendation:

- Plot the value of the function you are trying to minimise vs. iteration number
- Make sure it is decreasing
- Reduce eta if necessary

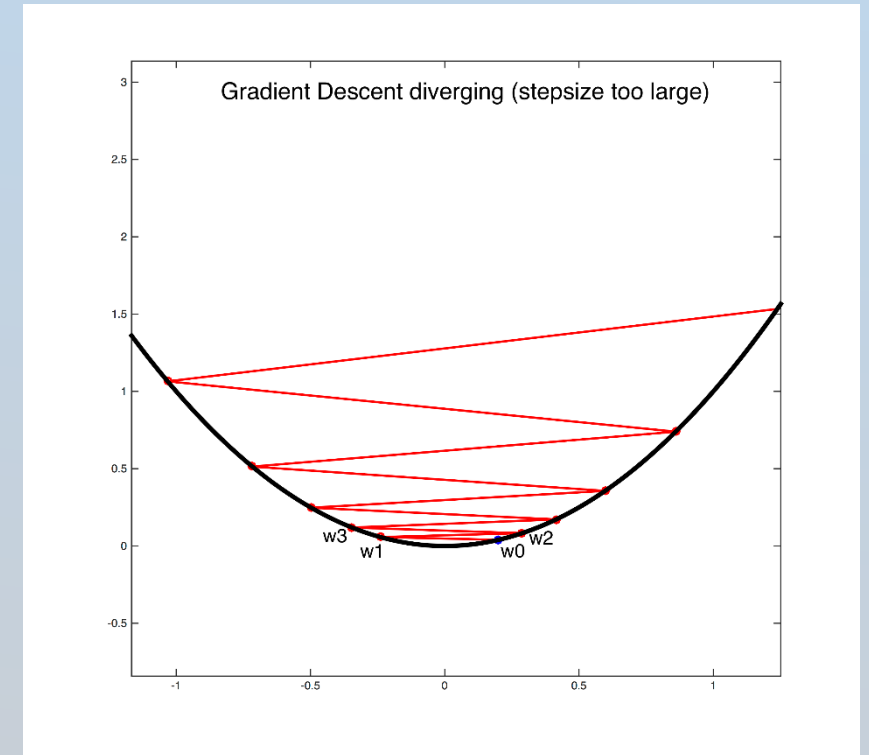


Image source: http://www.cs.cornell.edu/courses/cs4780/2017sp/lectures/images/gradient_descent_diverging.png

Example 1D Gradient Descent problem

Inefficiency:

- Recalculates the automatic differentiation formula every step of loop!

Fix this by:

1. pulling out guts of main loop into a separate python function
 - `def do_update():`
2. Annotate this function by “`@tf.function`”

Optimised version:

```
import tensorflow as tf
eta = 0.1 # learning rate
x = tf.Variable(10.0, tf.float32) # arbitrary initial value

@tf.function
def do_update(x):
    with tf.GradientTape() as tape:
        y=tf.pow(x,2.0)-4.0*x+4.0
    dydx=tape.gradient(y, x)
    x.assign(x-dydx*eta)
    return y

for i in range(50):
    y=do_update(x)
    print("iteration:",i, "x:", x.numpy(), "y:", y.numpy())
```

Put these changes into your notebook as a new script, under “Exercise1, Version 2”

Optimised version: @tf.function

- Use function annotation @tf.function to speed up execution, take advantage of the GPU, and for saving models
 - It allows tensorflow to cache the graph of computations so that it doesn't have to recalculate them (or the derivatives) every iteration.
- Adding the @tf.function in this task sped things up by around 4 times
- Put @tf.function around the main functionality of your training loop
 - After you've debugged things
 - Warning – the function you are optimising must not refer to any global variables (unless they are strictly constants)
- See [Better performance with tf.function | TensorFlow Core](#)

Using a built-in optimizer

- This will behave identically to our initial solution

```
import tensorflow as tf

eta = 0.1 # learning rate
x = tf.Variable(10.0, tf.float32) # arbitrary initial value

optimizer = tf.keras.optimizers.SGD(eta)
def calc_y():
    y=tf.pow(x,2.0)-4.0*x+4.0
    return y

for i in range(50):
    optimizer.minimize(calc_y, [x])
    print("iteration:",i, "x:", x.numpy(), "y:", calc_y().numpy())
```

Put these changes into your notebook as a new script, under “Exercise1 version 3”

Using a built-in optimizer

- Built-in optimizer

```
optimizer = tf.keras.optimizers.SGD(eta)
```

- SGD=Stochastic Gradient Descent
- Other optimizers are available
- How does it know what variables to optimise?

```
optimizer.minimize(calc_y, [x])
```

Exercise1, Version 3:

1. Modify your code to use `tf.keras.optimizers.SGD`
2. Check it behaves identically as before.

Using a built-in optimizer

```
optimizer = tf.keras.optimizers.SGD(eta)
```

- Other built-in optimizers work better with neural networks:

```
optimizer = tf.keras.optimizers.Adam()
```

```
optimizer = tf.keras.optimizers.RMSProp()
```

More on Gradient Descent

Gradient descent is heavily used in neural--
network optimisation.




Image source: http://www.cs.cornell.edu/courses/cs4780/2017sp/lectures/images/gradient_descent_diverging.png

Gradient Descent Local Minima

- Gradient descent will only head to the nearest local minimum




Image source: <https://static.thinkingandcomputing.com/2014/03/bprop.png>

- Might need several attempts from different random starting positions

Multidimensional Gradient Descent

- Gradient descent works too in multidimensional search spaces




Image source: <http://blog.datumbox.com/wp-content/uploads/2013/10/gradient-descent.png>

$$w_{t+1} = w_t - \eta \frac{\partial J}{\partial w}$$