

Introduction to TensorFlow and Deep Learning

Lecture 2: Neural Networks

IADS Summer School, 1st August 2022

Dr Michael Fairbank

University of Essex, UK

Email: m.fairbank@essex.ac.uk

Lecture Outline

- Summary so far:
 - TensorFlow basics
 - AutoDiff
 - Gradient Descent to solve a quadratic minimisation problem
- This Lecture (11.00am-12.30pm):
 - Feedforward neural networks
 - Training a network
 - (Using XOR benchmark problem)
 - Tricks for improving network training
 - Weight initialisation
 - Input data preparation
 - Choice of Activation Function
 - One-hot encoding and Cross Entropy loss function

Deep learning vs. Neural Networks

- Deep Learning = deep neural networks
 - Deep learning can be considered a re-branding of neural networks
 - Rebranded since neural networks started working much better
- Reasons for neural networks improving:
 - More training data
 - More processing power
 - Clever neural-network architectures
 - CNNs, Dropout, LSTM nodes, Transformers
 - Better activation functions and weight initialisations
 - ReLu. Xavier initialisation
 - Better learning algorithms (e.g. Adam)

What is a neural network?

A 3-4-2 feedforward neural network:

Takes in 3 input values


- a vector of length 3

Processes it

- hidden nodes “activate” with variable intensity
- output nodes “activate” with variable intensity
- Arrows are “weights” which amplify or reduce signals

Produces an output vector

- a vector of length 2 here



What is a neural network?

A feedforward neural network:

$$y = f(W, x)$$


- W =all “weights”, e.g. a list of tensors
- Give it enough weights, each appropriately set, and the neural network can represent any function.
 - “Universal function approximator”
 - E.g. could be a vision system
 - input=grid of pixel intensities of image
 - Output = classification of what the image is (car, house, etc)
 - Could be a language translation system
 - Input = vector of words in an English sentence
 - Output = vector of words in an French sentence

Feedforward algorithm

Number the nodes.

Label the weights:

w_{ij} connects node i to node j



Feedforward algorithm

The input vector to this network must be of dimension 3

- Input vector $\vec{x} = (x_1, x_2, x_3)$

The “electrical” signals passing along the wires to node 4 therefore total to:

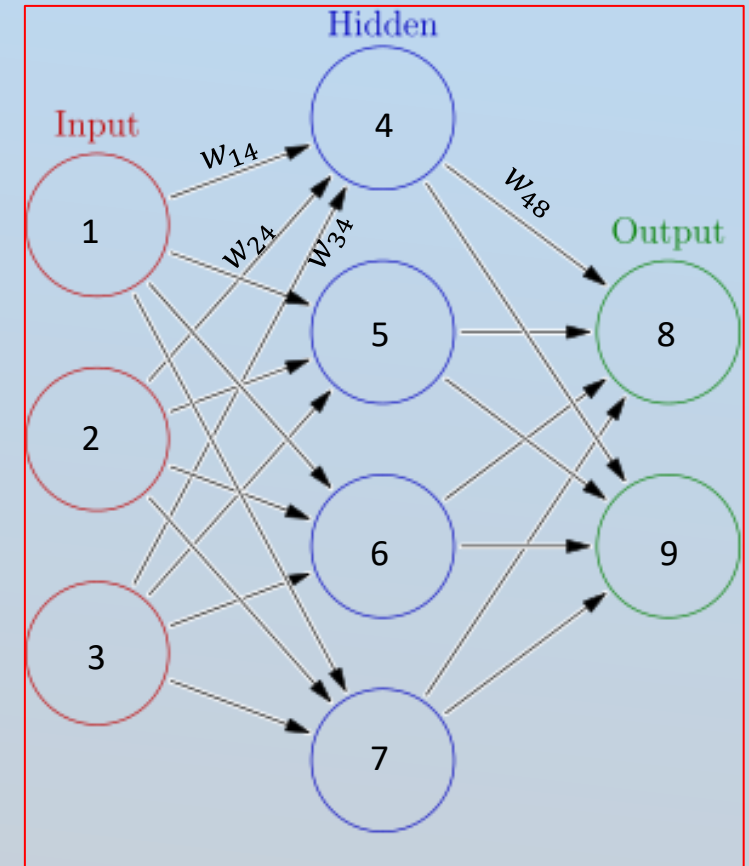
$$s_4 = x_1 w_{14} + x_2 w_{24} + x_3 w_{34}$$

Traditionally, node 4 will “fire” if s_4 is bigger than some threshold, b_4 , then

“Activation of node 4” $\longrightarrow a_4 = \begin{cases} 1 & \text{if } s_4 > b_4 \\ 0 & \text{if } s_4 \leq b_4 \end{cases}$

Equivalently, $a_4 = g(x_1 w_{14} + x_2 w_{24} + x_3 w_{34} - b_4)$

where $g(x) = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases}$ is called the “activation function” or “nonlinearity”

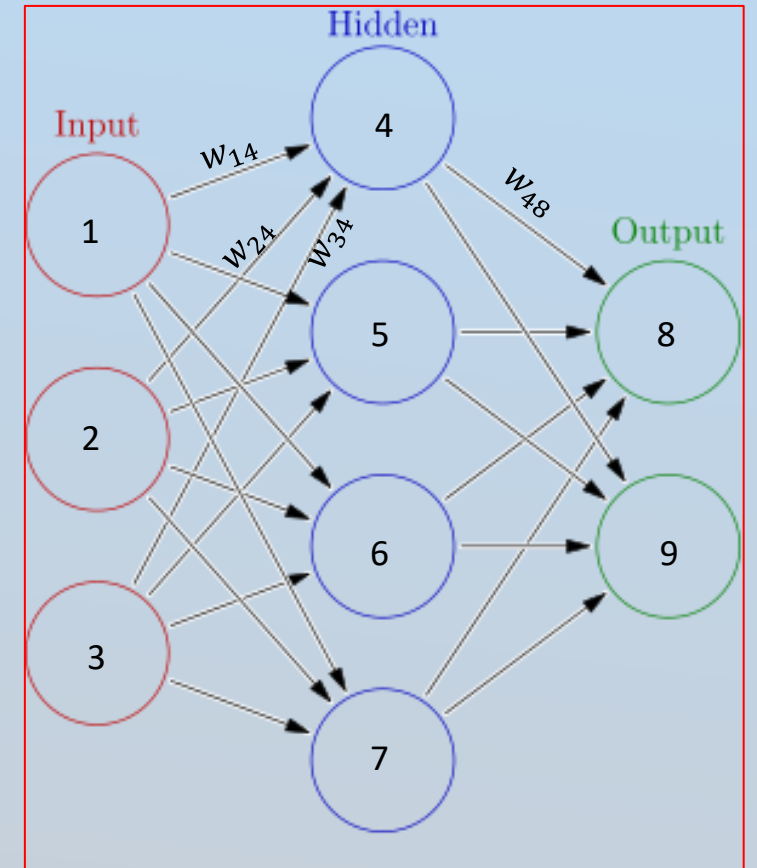


Feedforward algorithm

$$\begin{aligned}a_4 &= g(x_1w_{14} + x_2w_{24} + x_3w_{34} - b_4) \\a_5 &= g(x_1w_{15} + x_2w_{25} + x_3w_{35} - b_5) \\a_6 &= g(x_1w_{16} + x_2w_{26} + x_3w_{36} - b_6) \\a_7 &= g(x_1w_{17} + x_2w_{27} + x_3w_{37} - b_7)\end{aligned}$$

More concisely:


$$(a_4 \ a_5 \ a_6 \ a_7) = g\left((x_1 \ x_2 \ x_3) \begin{pmatrix} w_{14} & w_{15} & w_{16} & w_{17} \\ w_{24} & w_{25} & w_{26} & w_{27} \\ w_{34} & w_{35} & w_{36} & w_{37} \end{pmatrix} - (b_4 \ b_5 \ b_6 \ b_7)\right)$$



Feedforward algorithm

And for the output layer :

$$(a_8 \ a_9) = g \left((a_4 \ a_5 \ a_6 \ a_7) \begin{pmatrix} w_{48} & w_{49} \\ w_{58} & w_{59} \\ w_{68} & w_{69} \\ w_{78} & w_{79} \end{pmatrix} - (b_8 \ b_9) \right)$$




Feedforward algorithm

Together:

$$(a_4 \ a_5 \ a_6 \ a_7) = g \left((x_1 \ x_2 \ x_3) \begin{pmatrix} w_{14} & w_{15} & w_{16} & w_{17} \\ w_{24} & w_{25} & w_{26} & w_{27} \\ w_{34} & w_{35} & w_{36} & w_{37} \end{pmatrix} - (b_4 \ b_5 \ b_6 \ b_7) \right)$$

$$(a_8 \ a_9) = g \left((a_4 \ a_5 \ a_6 \ a_7) \begin{pmatrix} w_{48} & w_{49} \\ w_{58} & w_{59} \\ w_{68} & w_{69} \\ w_{78} & w_{79} \end{pmatrix} - (b_8 \ b_9) \right)$$



By “Training” the neural network, we aim to find values of all the w_{ij} and b_i values so that the neural network behaves as we want it to.

Feedforward algorithm

More concisely:

Hidden layer: $\vec{h1} = g(\vec{x}W1 + \vec{b1})$, where

$$\vec{h1} = (a_4 \quad a_5 \quad a_6 \quad a_7)$$

$$W1 = \begin{pmatrix} w_{14} & w_{15} & w_{16} & w_{17} \\ w_{24} & w_{25} & w_{26} & w_{27} \\ w_{34} & w_{35} & w_{36} & w_{37} \end{pmatrix}$$

$$\vec{x} = (x_1 \quad x_2 \quad x_3)$$


Note: moved minus sign

$$\vec{b1} = -(b_4 \quad b_5 \quad b_6 \quad b_7)$$

Output layer: $\vec{y} = g(\vec{h1}W2 + \vec{b2})$, where

$$\vec{y} = (a_8 \quad a_9) \quad W2 = \begin{pmatrix} w_{48} & w_{49} \\ w_{58} & w_{59} \\ w_{68} & w_{69} \\ w_{78} & w_{79} \end{pmatrix}$$


$$\vec{b2} = -(b_8 \quad b_9)$$



Feedforward algorithm


Reminder: The activation function is

$$g(x) = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases}$$


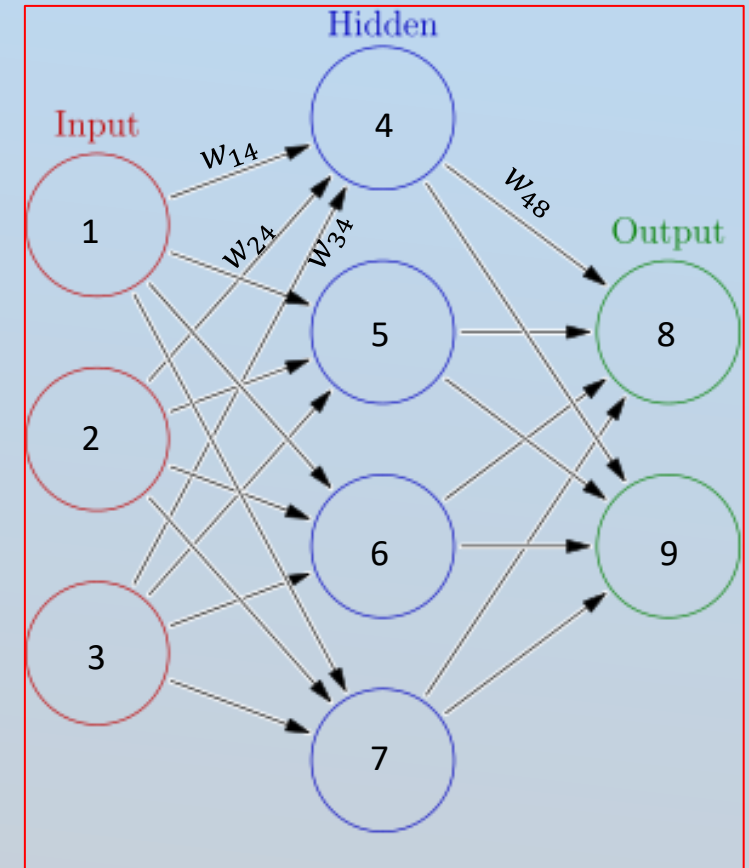


This is not suited for gradient descent (not differentiable)

Hence we change it to the logistic sigmoid function:



or any other non-linear, increasing, smooth function (e.g. $\tanh(x)$).



Feedforward algorithm

Summary, in TensorFlow code


For input row vector x :

```
h1=tf.sigmoid(tf.matmul(x,W1) + b1)
y=tf.sigmoid(tf.matmul(h1,W2) + b2)
```

A neural network coded in just 2 lines!

If we wanted a second hidden layer:

```
h1=tf.sigmoid(tf.matmul(x,W1) + b1)
h2=tf.sigmoid(tf.matmul(h1,W2) + b2)
y=tf.sigmoid(tf.matmul(h2,W3) + b3)
```



Note: for `tf.matmul(...)` to work, both arguments must be equal rank.

Therefore x must be rank 2

E.g. x has shape $[1, m]$ (i.e. a rank 2 row vector)

m =number of input nodes to network

Feedforward algorithm

Summary, in TensorFlow code

For input row vector x :

```
h1=tf.sigmoid(tf.matmul(x,W1) + b1)
y=tf.sigmoid(tf.matmul(h1,W2) + b2)
```

For the network shown, x has shape $[1,3]$

What shape does $W1$ have?

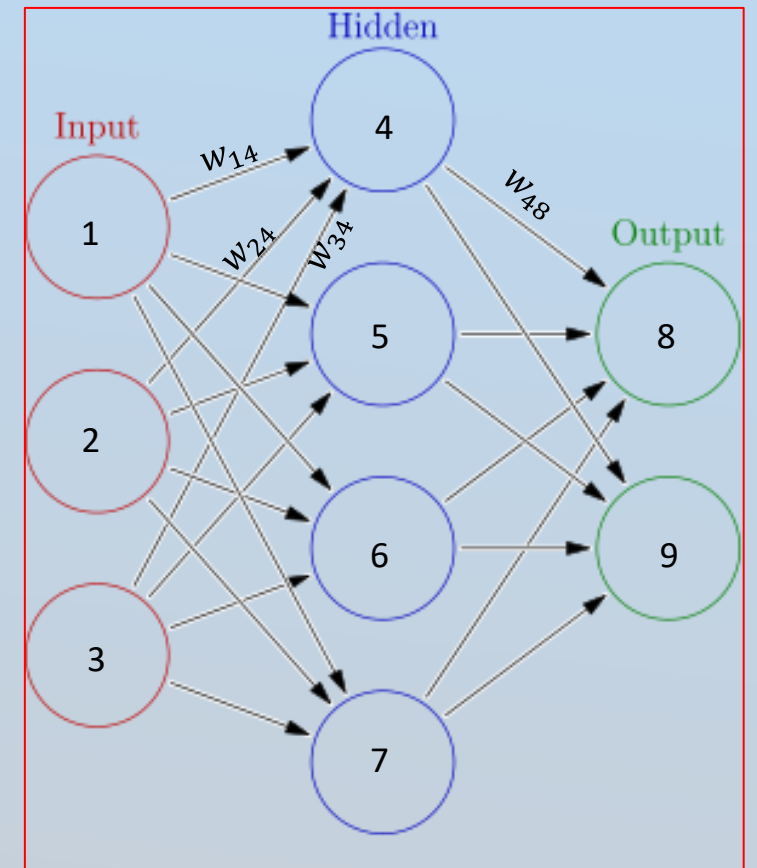
What shape does $W2$ have?

What shape does $b1$ have?

What shape does $b2$ have?

What shape does $h1$ have?


What shape does y have?



$x:[1,3]$ $b1:[1,4]$ $b2:[1,2]$
 $W1:[3,4]$ $W2:[4,2]$
 $h1:[1,4]$ $y:[1,2]$

Feedforward algorithm

The matrices W_1 , W_2 are called the “weight matrices” or “kernels”.



Feedforward algorithm

Exercise 1: (Template code is on next slide and in lecture2-notebook-ffnn)

1. Create a 2-2-1 feedforward network. (One hidden layer, of just 2 nodes).
2. Use sigmoid activation functions at each layer. `tf.sigmoid(...)`
3. Randomise weights and biases.

E.g. to set W1 to a random 4 by 5 matrix, use

```
W1=tf.Variable(tf.random.truncated_normal([4,5], stddev=0.1))
```

However, to ensure we all get the same results, we will set the random seed.

```
E.g. W1=tf.Variable(tf.random.truncated_normal([4,5], stddev=0.1, seed=1))
```

4. Then evaluate your network with input vector $\vec{x} = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$. Check this gives an output of $\vec{y} = \begin{bmatrix} 0.50787264 \end{bmatrix}$

(assuming all of our machines and python + TensorFlow versions have consistent random number generator algorithms)

Feedforward algorithm

Template script for exercise. Replace green # signs.

```
import tensorflow as tf

x=tf.constant([[1,1]], tf.float32) # our input vector

# Build our random weight and bias matrices, of appropriate shapes
W1 = tf.Variable(tf.random.truncated_normal([#, #], stddev=0.1, seed=1))
b1 = tf.Variable(tf.random.truncated_normal([#, #], stddev=0.1, seed=2))
W2 = tf.Variable(tf.random.truncated_normal([#, #], stddev=0.1, seed=3))
b2 = tf.Variable(tf.random.truncated_normal([#, #], stddev=0.1, seed=4))

# define our feed-forward neural network here:
def run_network(x):
    h1=#TODO
    y=#TODO
    return y

print(run_network(x).numpy())
```

This script is in lecture2-notebook-ffnn.ipynb under “Exercise 1”. Complete it now.

Randomising initial weights

- We need to randomise weights in a neural network before training
 - Symmetry breaking for gradient descent
 - Randomised start point for gradient descent




Image source: <https://static.thinkingandcomputing.com/2014/03/bprop.png>

Randomising initial weights

- We normally wouldn't set the random seed like we did
 - might need multiple starts from different random weights, to try to avoid local minima
- The magnitude you choose for the initial random weights is very important
 - This is one of the big breakthroughs in deep learning. (More later...)

Training the network

- To train the network we need to choose the weights and bias tensors so as to make the network behave as desired.
- Let's make it learn the XOR function.

4 input vectors

XOR truth table

Input		Output
A	B	
0	0	0
0	1	1
1	0	1
1	1	0

4 target output vectors

The diagram illustrates the XOR truth table with four rows of data. Four blue arrows point from the text '4 input vectors' to the first two columns (A and B) of the table. Another four blue arrows point from the text '4 target output vectors' to the third column (Output) of the table.

Training the network

- We can modify our feedforward network to push through all 4 input vectors at once:
- Change `x=tf.constant([[1,1]], tf.float32)`
To `x=tf.constant([[0,0],[0,1],[1,0],[1,1]], tf.float32)`
- The equations still work fine:
`h1=tf.sigmoid(tf.matmul(x,W1) + b1)`
`y=tf.sigmoid(tf.matmul(h1,W2) + b2)`
- **Exercise 2:** Try this modification to your program.
 - Use the jupyter notebook template under “Exercise 2”
- Broadcasting is used for the “+b1” and “+b2”
- Our output tensor will now have shape [4,1]

XOR truth table

Input		Output
A	B	
0	0	0
0	1	1
1	0	1
1	1	0

Training the network

- Our network outputs $\vec{y} = \begin{pmatrix} 0.509233 \\ 0.509222 \\ 0.507883 \\ 0.507872 \end{pmatrix}$. However our targets were $\vec{t} = \begin{pmatrix} 0 \\ 1 \\ 1 \\ 0 \end{pmatrix}$

Assuming our RNG seeds behave the same.

Note: Non-deterministic behaviour in Jupyter; need to restart jupyter kernel to reset tensorflow's random state.

- The squared error between these is
 - $L = \|\vec{y} - \vec{t}\|^2$
- We train the network by doing gradient descent on this

$$w \leftarrow w - \eta \frac{\partial L}{\partial w}$$

Where w here represents *all* weights and biases

$$w = (W1, W2, b1, b2)$$

Training the network

$$L = \|\vec{y} - \vec{t}\|^2$$

$$w \leftarrow w - \eta \frac{\partial L}{\partial w}$$

Notes:

1. L is our “loss function”, or “error function” or “cost function”
2. Tensorflow will calculate these partial derivatives $\frac{\partial L}{\partial w}$ for us using Autodiff (backpropagation)
3. Tensorflow will implement gradient descent $w \leftarrow w - \eta \frac{\partial L}{\partial w}$ using its canned optimizer
4. First we need to define L in terms of tensorflow code

Training the network

$$L = \|\vec{y} - \vec{t}\|^2$$

Define L in terms of TensorFlow code:

$\|A\|^2$ means sum and square every component of A

$$L = \|\vec{y} - \vec{t}\|^2$$

```
deltas=tf.subtract(y, y_labels)
squared_deltas=tf.square(deltas)
loss=tf.reduce_sum(squared_deltas)
```

Q: Why is this function called “reduce”...?

Exercise 3: Train your network

1. Define your constant 4*1 tensor *y_labels*
 - Make sure it is shape [4,1] not shape [1,4]
 - Think about the nesting of your square brackets to achieve this.
2. Define your loss function

```
def calc_loss():  
    y=run_network(x)  
    deltas=tf.subtract(y, y_labels)  
    squared_deltas=tf.square(deltas)  
    loss=tf.reduce_sum(squared_deltas)  
    return loss
```

XOR truth table

Input		Output
A	B	
0	0	0
0	1	1
1	0	1
1	1	0

Exercise 3: Train your network

3. Train your network in a loop:

```
optimizer = tf.keras.optimizers.SGD(0.5)
for i in range(20000):
    optimizer.minimize(calc_loss, [W1,b1,W2,b2])
    if (i%1000)==0:
        print("iteration ",i," loss", calc_loss().numpy())
print(run_network(x).numpy())
```

XOR truth table

Input		Output
A	B	
0	0	0
0	1	1
1	0	1
1	1	0

Make your changes in the jupyter notebook under “Exercise 3”

Gotchas from this exercise:

Check targets and y are the same shape, i.e. both 4×1 and not 1×4 .

(If you get this wrong then “deltas” will acquire size 4×4 (by broadcasting), and reduce_sum will give a strange result... and during training “loss” will never go below 4)

Exercise 3: Train your network

Review: Train your network

- Note that these 3 lines

```
optimizer = tf.keras.optimizers.SGD(eta)
for i in range(20000):
    optimizer.minimize(calc_loss, [W1,b1,W2,b2])
```

- were shorthand for:

```
for i in range(20000):
    with tf.GradientTape() as tape:
        L=calc_loss()
        [dLdW1,dLdW2,dLdb1,dLdb2]=tape.gradient(L, [W1,W2,b1,b2])
        W1.assign(W1-dLdW1*eta)
        W2.assign(W2-dLdW2*eta)
        b1.assign(b1-dLdb1*eta)
        b2.assign(b2-dLdb2*eta)
```

- ...and the “tape.gradient” line is itself a massive shorthand
 - for the backpropagation derivative calculation...

XOR + Universal Function approximation

We've solved XOR. This is a tiny network. Modern deep learning networks are tens or hundreds of layers deep.

XOR relates to the universal function approximation capabilities of neural networks.

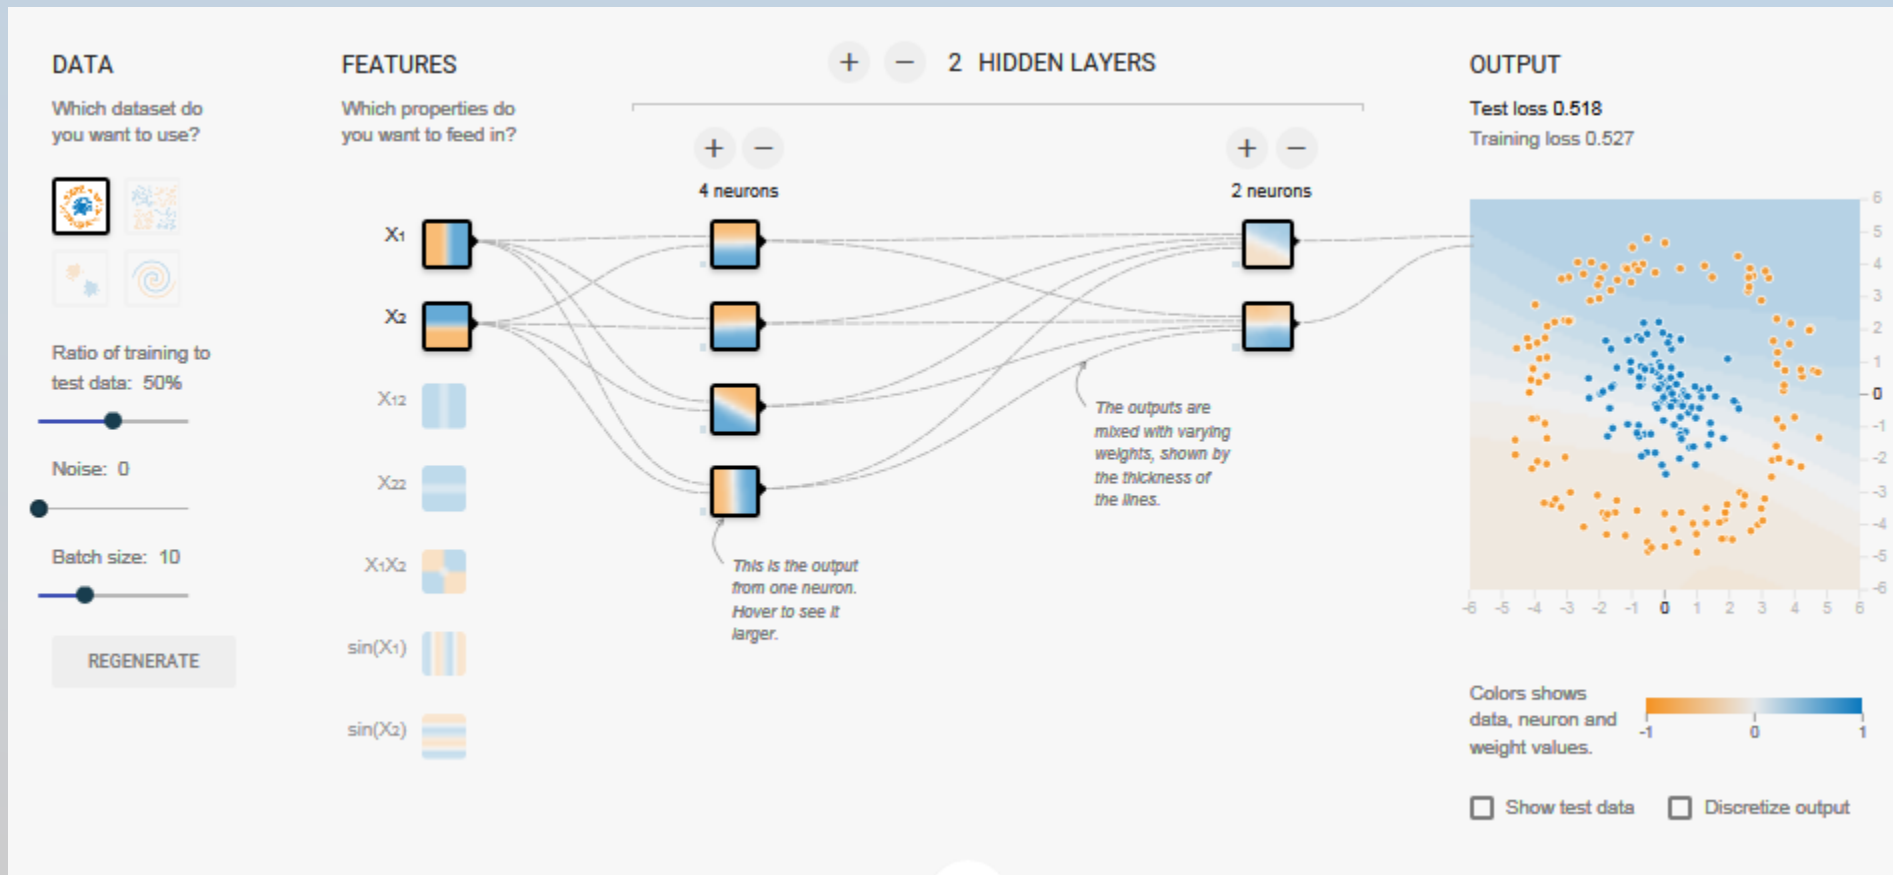
XOR truth table

Input		Output
A	B	
0	0	0
0	1	1
1	0	1
1	1	0

TensorFlow playground website

Interesting website for building intuition of what a neural network is:

<https://playground.tensorflow.org/>



Running the neural network on new data

- A “pattern” is a pairing of input and output vectors
 - Our network has learned 4 “patterns”.
- What if we wanted to input a new pattern? E.g. input (0.5,0.5)
 - How would we rewrite our TensorFlow code?

XOR truth table

Input		Output
A	B	
0	0	0
0	1	1
1	0	1
1	1	0

Running the neural network on new data

Exercise: try this.

Then we can evaluate our network on a new pattern, e.g.

```
print(run_network ( [[0.5,0.5]]).numpy())
```

XOR truth table

Input		Output
A	B	
0	0	0
0	1	1
1	0	1
1	1	0

Feedforward neural nets: Higher level functions

To create two network layers, instead of:

```
W1 = tf.Variable(tf.random.truncated_normal([2,2], stddev=0.11))
b1 = tf.Variable(tf.random.truncated_normal([1,2], stddev=0.1))
W2 = tf.Variable(tf.random.truncated_normal([2,1], stddev=0.1))
b2 = tf.Variable(tf.random.truncated_normal([1,1], stddev=0.1))
def run_network(x):
    h1=tf.sigmoid(tf.matmul(x,W1) + b1)
    y=tf.sigmoid(tf.matmul(h1,W2) + b2)
    return y
```

- “layer1” and “layer2” act as callable functions
- You can access the weights and biases for a layer as
 - “layer1.trainable_variables”, which evaluates as [W1,b1]
 - Note: You must run the network with an input tensor x before accessing trainable_variables though

We can just do:

```
layer1=tf.keras.layers.Dense(2, activation=tf.nn.sigmoid)
layer2=tf.keras.layers.Dense(1, activation=tf.nn.sigmoid)
def run_network(x):
    h1=layer1(x)
    y=layer2(h1)
    return y
```

This will create all 4 Variable weight and bias tensors for us, and randomise them appropriately

Feedforward neural nets: Higher level functions

Furthermore, instead of:

```
layer1=tf.keras.layers.Dense(2, activation=tf.nn.sigmoid)
layer2=tf.keras.layers.Dense(1, activation=tf.nn.sigmoid)
def run_network(x):
    h1=layer1(x)
    y=layer2(h1)
    return y
```

We can just do:

```
layer1=tf.keras.layers.Dense(2, activation=tf.nn.sigmoid)
layer2=tf.keras.layers.Dense(1, activation=tf.nn.sigmoid)
model = tf.keras.Sequential([layer1,layer2])
def run_network(x):
    return model(x)
```

`model.trainable_variables` will give us
`[W1,b1,W2,b2]`

Exercise 4 (Optional): rewrite your code to use one of these 2 methods.

Limitations of XOR

To get deeper networks trained, we need
better weight initialisation,
better activation functions,
more training data
more CPU time.

Challenge: Enhance this code

Try to put the key line which computes gradients, i.e.,

```
optimizer = tf.keras.optimizers.SGD(0.5)
for i in range(20000):
    optimizer.minimize(calc_loss, [W1,b1,W2,b2])
    if (i%1000)==0:
        print("iteration ",i," loss", calc_loss().numpy())
print(run_network(x).numpy())
```

Into a separate python function which has the @tf.function annotation

Q: Why do this?

Tricks for improving network training

Input pre-processing

Neural networks work best when input vectors are normalised to have variance 1.

So for each real-valued column of your dataset input tensor, subtract off column means...

```
import numpy as np
inputs=[[0.0,0],[0,10],[10,0],[10,1]]
column_means=np.mean(inputs, axis=0)
inputs-=column_means
```

and then divide each column by column standard deviation.


```
column_stds=np.std(inputs, axis=0)
inputs/=column_stds
```

Don't forget to do the same to any test data you later feed the neural network
(carefully record the above column_means and column_stds for this purpose)

For more sophisticated pre-processing, see [whitening](#)

This slide is really important:


- Normalise your inputs.
- Beginners forget this step




Better Activation functions

In increasing order of usefulness for deep learning:

1. Logistic Sigmoid function. $g(x) = \frac{1}{1+e^{-x}}$. `tf.sigmoid(X)`




2. Tanh function. $g(x) = \tanh(x)$. `tf.tanh(X)`



3. ReLu (Rectified Linear)function.

$$g(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases} \cdot \text{tf.nn.relu(X)}$$



ReLU is now the most commonly used activation function for training deep neural networks. It has nice properties in that learning gradients are less likely to decay / explode.

See [https://en.wikipedia.org/wiki/Rectifier_\(neural_networks\)](https://en.wikipedia.org/wiki/Rectifier_(neural_networks))

https://en.wikipedia.org/wiki/Vanishing_gradient_problem

https://www.tensorflow.org/api_docs/python/tf/nn/relu

Better Weight Initialisation

We need random initialised weights to help gradient descent, and break symmetry. The magnitudes of those initial random weights is important, to try to prevent vanishing / exploding learning gradients.

Glorot/Xavier initialisation: Set the standard-deviation of for a weight matrix of dimension $n \times m$

$$\sigma = \sqrt{\frac{2}{n + m}}$$

The above formula varies depending on the choice of:

- Activation function used
- Whether you used normal distribution or Gaussian randomisation
- The exact research you follow

Source:

Glorot, Xavier, and Yoshua Bengio. "Understanding the difficulty of training deep feedforward neural networks." Proceedings of the Thirteenth International Conference on Artificial Intelligence and Statistics. 2010.

<http://www.jmlr.org/proceedings/papers/v9/glorot10a.html>

Better Weight Initialisation

Glorot/Xavier initialisation: $\sigma = \sqrt{\frac{2}{n+m}}$

Glorot/Xavier initialisation is easy when using the higher level functions to create your network layers:

```
layer1=tf.keras.layers.Dense(2, activation=tf.nn.sigmoid, kernel_initializer='glorot_uniform')
```

In fact Glorot/Xavier initialisation is default here.

See https://www.tensorflow.org/api_docs/python/tf/keras/initializers/ for further information

Categorical inputs + one-hot encoding

If we have an input to a NN that represents a category, e.g. “profession:”

1=Teacher

2=Doctor

3=Student

4=Unemployed

Since there is no logical ordering of these 4 items numerically, we should not give them as a *single numerical input* to the NN.

Instead use *four different inputs* to the NN, with “one-hot encoding”:

(1,0,0,0)=Teacher

(0,1,0,0)=Doctor

(0,0,1,0)=Student

(0,0,0,1)=Unemployed

Categorical inputs + one-hot encoding

Q. How many inputs should we use for “age of child (1-10)”?

Q. How many inputs should we use for “colour of eyes”?

Q. How many inputs should we use for “position of tennis ball”?

Q. How many outputs should we have for a letter-of-alphabet reading (A-Z) vision system?

Similarly, target outputs corresponding to different categories should be one-hot encoded

You can encode into one-hot encoded via `tf.one_hot` function

https://www.tensorflow.org/api_docs/python/tf/one_hot

Categorical inputs + one-hot encoding

Q: what is

```
tf.argmax(tf.constant([[0,0,0,0,1], [0,0,0,1,0]]),axis=1)
```

If our output tensor is y (one-hot encoded) then `tf.argmax(y, axis=1)` will reverse the one-hot encoding.

If the targets are `y_targets` (integer class number; so NOT one-hot encoded), then we can calculate the number of correct rows as follows:

```
correct_prediction = tf.equal(tf.argmax(y, axis=1), y_targets)  
corrent_count = tf.reduce_sum(tf.cast(correct_prediction, tf.float32))
```

Training Classification Networks

Training Classification Networks

For classification systems, e.g. a vision system, it is usual to try to make the neural network output “probabilities” for each classification possibility.

Suppose we have possible 4 furniture categories: “Chair/Table/Wardrobe/Door”.

Hence we need an NN with 4 outputs.

Suppose the neural output vector is $\vec{y} = (-0.2, 0.5, 0.1, 0.1)$

Which furniture item do you think this means?

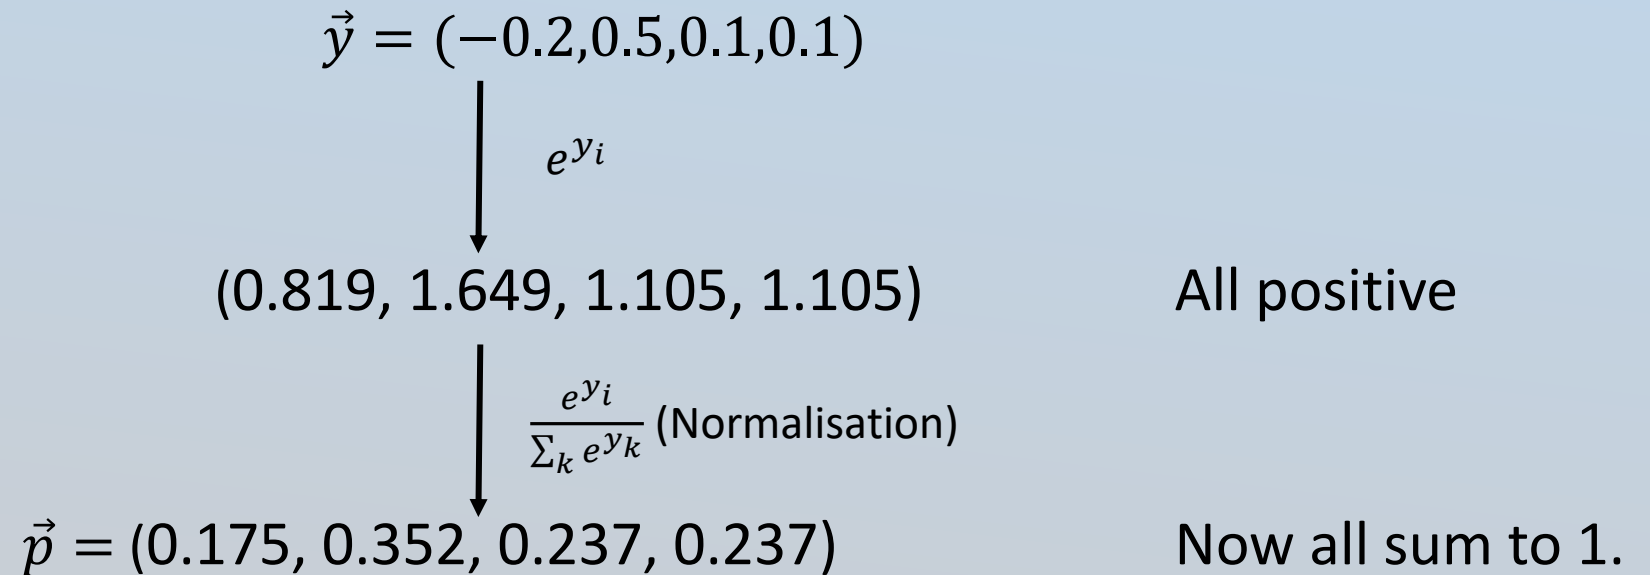
How confident are you?

How can we convert this NN output into “probabilities”?

Probabilities must all be positive and must sum to 1

Training Classification Networks

Probabilities must all be positive and must sum to 1



Summary: $p_i = \frac{e^{y_i}}{\sum_k e^{y_k}}$ (Softmax function)

Q: Does $P(\text{Table})=0.352$?

Cross-entropy Loss function

Now the neural network outputs produce \vec{p} with $p_i = \frac{e^{y_i}}{\sum_k e^{y_k}}$ (Softmax function)

We want these to eventually be trained to become true probabilities.

To achieve this, we train the NN with “Cross-Entropy Loss Function”

The cross-entropy loss for a single pattern is defined to be

$$L = - \sum_i t_i \log(p_i)$$

where \vec{t} is the one-hot encoded true output classification, and i is the category index.

Cross-entropy Loss function

The cross-entropy loss for a single pattern is defined to be

$$L = -\sum_i t_i \log(p_i)$$

where \vec{t} is the one-hot encoded true output classification, and i is the category index.

E.g. Consider a “table”, $\vec{t}=(0,1,0,0)$, with actual neural-network output $\vec{y} = (-0.2,0.5,0.1,0.1)$

The category index of “table” is 1.

Hence $t_1 = 1$ $t_0 = t_2 = t_3 = 0$.

Hence $L = -\log(p_1)$

In this example, when training attempts to minimise L , it will have to minimise $L = -\log(p_1)$, i.e. maximise p_1 , i.e. maximise y_1 compared to all of y_0, y_2, y_3

Cross-entropy Loss function

The cross-entropy loss for a single pattern is defined to be

$$L = - \sum_i t_i \log(p_i)$$

where \vec{t} is the one-hot encoded true output classification, and i is the category index.

Where does this formula come from? It can be shown that minimising the above L over your whole dataset is the same as *finding the weights which maximise your neural network's estimation of the total probability that your given training dataset members each have their given (fixed) individual category label.*

Example: If you trained it with 60 images of dogs and 40 images of cats, and your neural network could memorise all 100 images exactly, then it should learn to output $P(dog) \approx 1$ for every dog image and $P(cat) \approx 1$ for every single cat image.

- If the neural network is then given a new image of a dog we *hope* that the image is sufficiently similar to the other dog images such that the neural network gives $P(dog) \approx 1$ for that new dog image. But this is not guaranteed of course.

Example 2: if you had 100 identical photos in your training dataset and 40% of them happened to be labelled “1” and the other 60% were labelled “2”, then the neural network would have to converge to giving a probability of 40% to category “1”.

- In this bizarre example, there is nothing the neural network can do to distinguish category 1 from category 2 (since the training input photos were all identical).

Cross-entropy Loss function

SoftMax in TensorFlow:

```
y=tf.keras.activations.softmax(y, axis=1)
```

Cross entropy in TensorFlow:

```
cce=tf.keras.losses.SparseCategoricalCrossentropy()  
cross_entropy = tf.reduce_mean(cce(y_true=train_labels, y_pred=y))
```

- This expects `y_pred` to have had softmax applied to it on entry.
- Also it expects `train_labels` to be an integer for the category number (so the label is NOT one-hot encoded)

The variable `cross_entropy` is returned by our “`calc_loss()`” function, to optimise through gradient descent

```
optimizer.minimize(calc_loss, trainable_variables)
```

Further Reading (Cross Entropy Loss)

Further reading:

[Where did the Binary Cross-Entropy Loss Function come from? | by Rafay Khan | Towards Data Science](#)

Also on cross entropy: [The Most Important Function in Machine Learning](#) (YouTube, P.Solver, 30mins)

Summary: Classification vs. Regression Problems

When you want a neural network to classify something-

- Use SoftMax to obtain probabilities
- Use cross entropy loss function to train
- Use max to pick the NN's favoured category

When you want a neural network to solve a real valued problem (A regression problem):

- Train with sum-of-squared error
- Don't use SoftMax on final layer

Optional Lecture Recap: How does a neural network work?

We have already studied “what a neural network is” now.

But a slick recap, especially on how data is propagated forwards through a network, is given by YouTuber Veritasium: [Future Computers Will Be Radically Different – YouTube](#)

Watch from 3m42s to 13m40s

0m00-3m42: Analogue vs digital computers

3m42-8m00: Single-layer neural networks (“Perceptron”)

8m00-10m00: Multi-layer neural networks

10m00-13m40s: Deep neural networks, Image Net